

# Pipelines as Code for Infrastructure at Scale

# Pipelines as Code

# **Pipelines as Code for Infrastructure**

# **Pipelines as Code for Infrastructure at Scale**

# **Agenda**

**A problem to solve**

**An attempt**

**The Issue**

**What are the options?**

**Programmatic pipelines**

**Scaling up**

**Results**

**Takeaways**

# Mario Fernandez




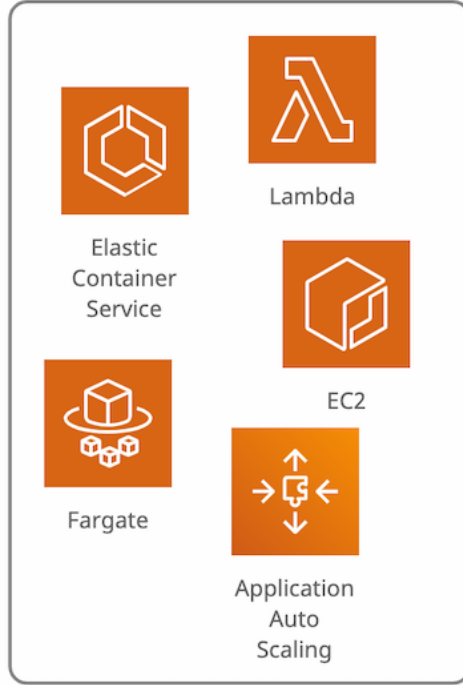
ThoughtWorks

**A problem to  
solve**


**Enable teams to move from on-premise to the  
cloud**





**Layer on top of AWS to provide a ready-to-use solution**









**An attempt**

**Provisioning a ton of infrastructure**



**Terraform**

## modules

- └─ **access**-from-other-accounts
- └─ application-environment-1
- └─ application-environment-2
- └─ cluster-{backup,management}
- └─ connected-hosted-zone
- └─ dispatcher-vpc
- └─ domain
- └─ functional-area
- └─ functional-area-{**access**,runtime}
- └─ logging
- └─ monitoring
- └─ operations
- └─ **private**-egress
- └─ **private**-ingress-{app,dispatcher}
- └─ system-services



```
modules/  
├─ access-from-other-accounts  
|   ├─ dependencies.tf  
|   ├─ locals.tf  
|   ├─ main.tf  
|   ├─ outputs.tf  
|   ├─ providers.tf  
|   └─ variables.tf  
├─ application-environment-1  
|   ├─ data.tf  
|   ├─ locals.tf  
|   ├─ main.tf  
|   ├─ outputs.tf  
|   ├─ providers.tf  
|   └─ variables.tf  
└─ tracing
```

```
[terragrunt] 2021/02/01 20:28:39 Executing hook: tflint
[terragrunt] 2021/02/01 20:28:39 Running command: tflint
[terragrunt] 2021/02/01 20:28:40 Running command: terraform apply
module.aws_vpc.module.vpc.aws_subnet.private[2]: Refreshing state..
module.aws_vpc.module.vpc.aws_vpc.this[0]: Refreshing state..
module.aws_vpc.module.vpc.aws_subnet.private[0]: Refreshing state..
module.aws_vpc.module.vpc.aws_subnet.private[1]: Refreshing state..
module.aws_vpc.module.vpc.aws_vpc.this[0]: Refreshing state..
module.aws_eks.data.aws_region.current: Refreshing state..

Apply complete! Resources: 0 added, 0 changed, 0 destroyed.
```

# **Spoiler Alert**

**Doing it by hand did not scale**



**[thoughtworks.com/radar/techniques/pipelines-for-infrastructure-as-code](https://thoughtworks.com/radar/techniques/pipelines-for-infrastructure-as-code)**

**Apply infrastructure changes automatically**

**Small modules**

**Pipeline is an orchestrator of existing scripts**

**The issue**





**YAML is really verbose**

**Like, really**

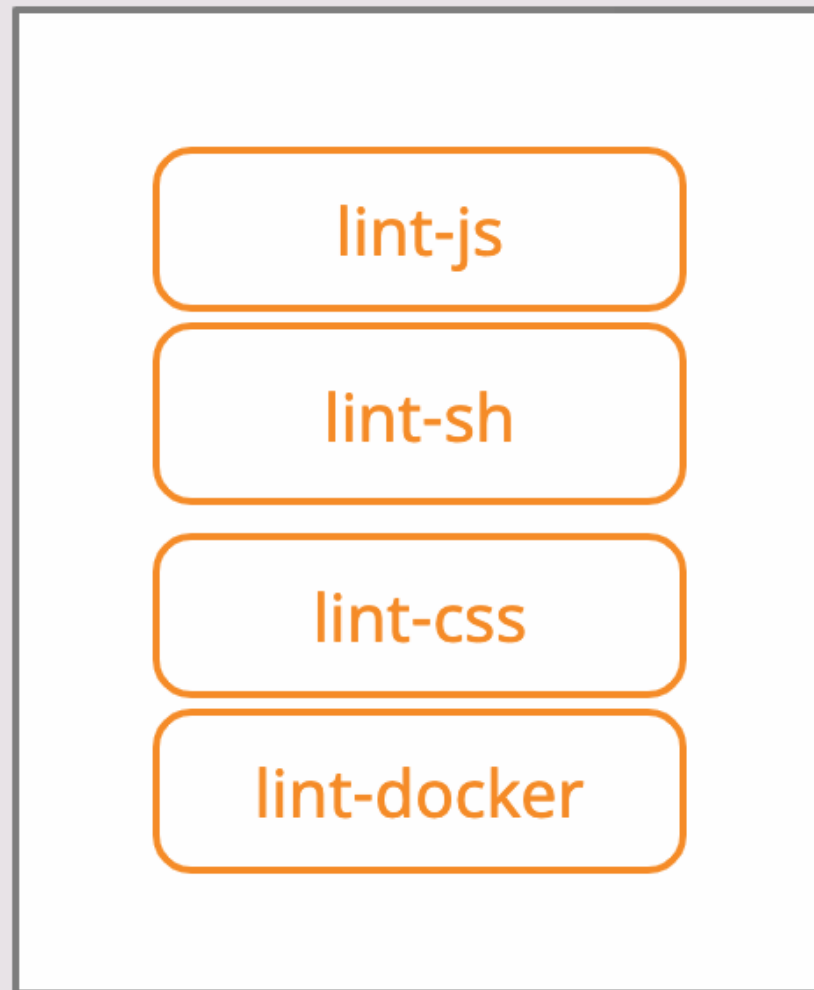
**Concourse doesn't help, either**

*Job*

*Inputs*



*Tasks*



```
- name: lint
  serial: true
  plan:
    - in_parallel:
        - get: git
          passed: [prepare]
          trigger: true
        - get: dev-container
          passed: [prepare]
    - in_parallel:
        - task: lint-sh
          image: dev-container
          params:
            <<: *common-params
            TARGET: sh
          file: git/pipeline/tasks/linter/task.yml
```

**Our ability to manage duplication is limited**

# **Anchor parameters**

```
- in_parallel:
  - task: lint-sh
    image: dev-container
    params:
      <<: *common-params
      TARGET: sh
    file: git/pipeline/tasks/linter/task.yml
```

```
common-params: &common-params
  CI: true
```



# **Parametrized tasks**

```
- in_parallel:
  - task: lint-sh
    image: dev-container
    params:
      <<: *common-params
      TARGET: sh
    file: git/pipeline/tasks/linter/task.yml
```

**YAML overdose!**

```
353 | serial: true
354 | - name: cluster-base-network-layout-test
355 |   plan:
356 |     - in_parallel:
357 |       - get: source-code
358 |         passed:
359 |           - cluster-global-dns
360 |           - functional-area-account-dev
361 |         trigger: true
362 |       - get: dev-tools
363 |         trigger: false
364 |       - get: terraform-tools
365 |         trigger: false
366 |       - get: container
367 |         trigger: false
368 |     - in_parallel:
369 |       - file: source-code/ci/tasks/provision.yaml
370 |         image: container
371 |         params:
372 |           AUTOMATION: true
373 |           CI_AWS_ACCESS_KEY_ID: ((aws-access-key-id))
374 |           CI_AWS_SECRET_ACCESS_KEY: ((aws-secret-access-key))
375 |           CONCOURSE_ADMIN_PASSWORD: ((admin-password))
376 |           CONCOURSE_ADMIN_USERNAME: ((admin-username))
377 |           REGION: eu-central-1
378 |         task: domain-eu-central-1
379 |     - in_parallel:
380 |       - file: source-code/ci/tasks/provision.yaml
381 |         image: container
382 |         params:
383 |           AUTOMATION: true
384 |           CI_AWS_ACCESS_KEY_ID: ((aws-access-key-id))
385 |           CI_AWS_SECRET_ACCESS_KEY: ((aws-secret-access-key))
```

**What are the  
options?**

**Something on top of YAML**

# **Three approaches**







## Templating in YAML

---

APR  
2019

**HOLD** ?

As infrastructures grow in complexity, so do the configuration files that define them. Tools such as **AWS CloudFormation**, **Kubernetes** and **Helm** expect configuration files in JSON or YAML syntax, presumably in an attempt to make them easy to write and process. However, in most cases, teams quickly reach the point where they have some parts that are similar but not quite the same, for example, when the same service must be deployed in different regions with a slightly different setup. For such cases tools offer **templating in YAML** (or JSON), which has caused a huge amount of **frustration with practitioners**. The problem is that the syntax of JSON and YAML requires all sorts of awkward compromises to graft templating features such as conditionals and loops into the files. We recommend using an API from a programming language instead or, when this is not an option, a templating system in a programming language, either a general-purpose language such as Python or something specialized such as **Jsonnet**.

■ We recommend using an API from a programming language, or, a templating system

# Programmatic pipelines

**Generate most of the pipeline automatically**

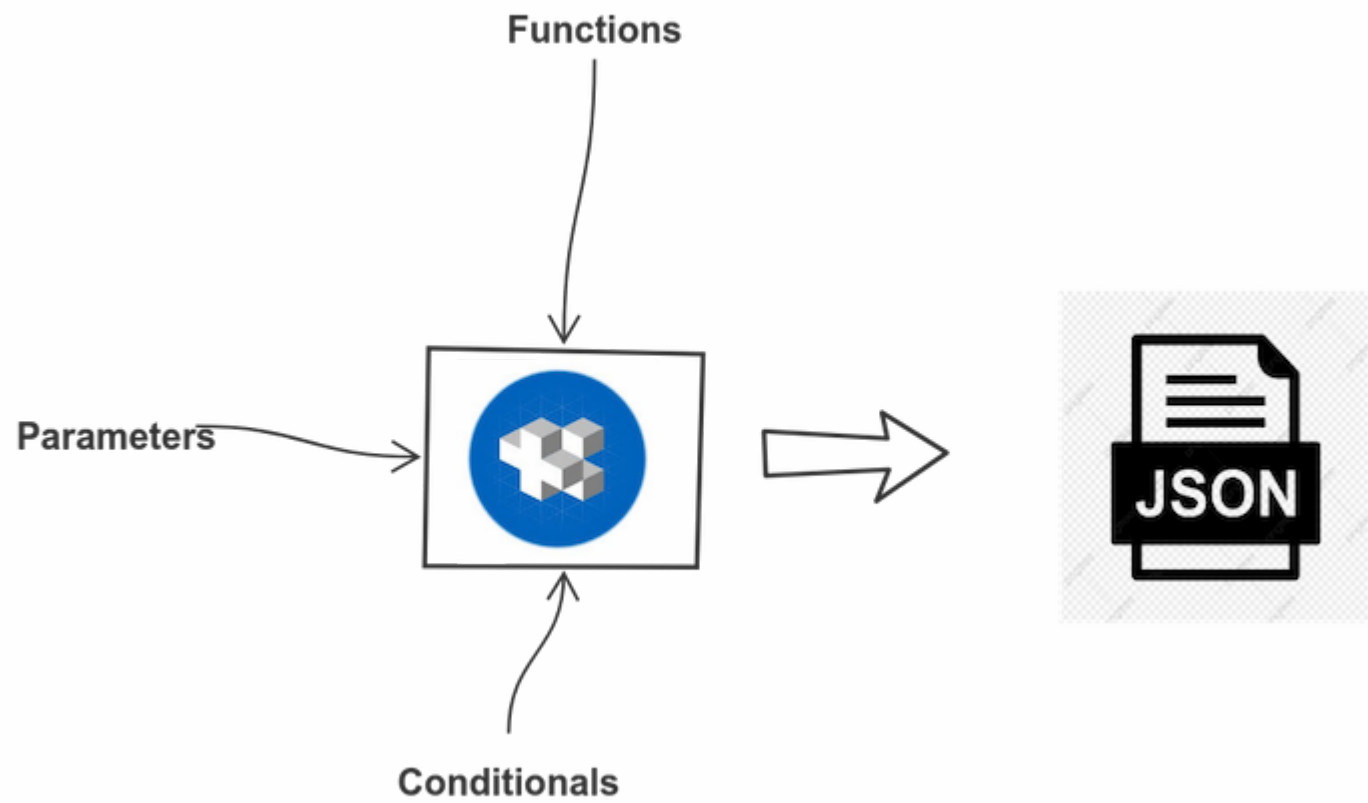
**Handle duplication through a parametrized approach**

# Jsonnet

**[jsonnet.org/](https://jsonnet.org/)**



**A data templating language for app and tool  
developers**



**YAML happens to be a superset of JSON**


# **Abstract building blocks**

*Job*

*Resource*



*Parallel*



```
local Parallel(tasks) = {  
  in_parallel: tasks  
}
```

```
local Job(name, serial = true, plan = []) = std.prune({  
  name: name,  
  serial: serial,  
  plan: plan  
})
```

```
DockerResource(name,  
                repository,  
                tag = 'latest', allow_insecure = false) = {  
  name: name,  
  type: 'docker-image',  
  source: {  
    repository: repository,  
    tag: tag  
  } + (  
    if allow_insecure then {  
      insecure_registries: [std.split(repository, '/')][0]]} else {}  
  ),  
}
```



```
local concourse = import 'concourse.libsonnet';  
{  
  docker: concourse.DockerResource('serverspec-container',  
                                     'sirech/dind-ruby', tag = '2.6.3'),  
}
```

```
{  
  "docker": {  
    "name": "serverspec-container",  
    "source": {  
      "repository": "sirech/dind-ruby",  
      "tag": "2.6.3"  
    },  
    "type": "docker-image"  
  }  
}
```

**[github.com/sirech/concourse-jsonnet-utils](https://github.com/sirech/concourse-jsonnet-utils)**

# Building your own DSL

**Reflect your conventions and idioms in code**

```
local source = 'git';  
local container = 'dev-container';  
  
local concourse = import 'concourse.libsonnet';  
  
local Inputs(dependencies = []) = concourse.Parallel(  
  [concourse.Get(s, dependencies = dependencies)  
    for s in [source, container]]  
);
```

```
local Task(name, file = name, image = container, params = {}) = {  
  task: name,  
  image: image,  
  params: {  
    CI_AWS_ACCESS_KEY_ID: ((aws-access-key-id)),  
    CI_AWS_SECRET_ACCESS_KEY: ((aws-secret-access-key)),  
    MODULE: module,  
    PRODUCT: product,  
    CI: true  
  } + params,  
  file: '%s/pipeline/tasks/%s/task.yml' % [source, file]  
};
```

**Back to the original example**

```
- name: lint
  serial: true
  plan:
    - in_parallel:
        - get: git
          passed: [prepare]
          trigger: true
        - get: dev-container
          passed: [prepare]
    - in_parallel:
        - task: lint-sh
          image: dev-container
          params:
            <<: *common-params
            TARGET: sh
          file: git/pipeline/tasks/linter/task.yml
```



```
local concourse = import 'concourse.libsonnet';

concourse.Job('lint', plan = [
  Inputs('prepare'),
  concourse.Parallel(
    [Task('lint-%s' % lang, 'linter', params = { TARGET: lang })
      for lang in ['sh', 'js', 'css', 'docker']]
  )
],
```

**[github.com/sirech/example-concourse-pipeline](https://github.com/sirech/example-concourse-pipeline)**

```

local product = std.extVar('SERVICE_NAME');
local config = std.extVar('CONFIG');

local concourse = import 'concourse.libsonnet';
local envs = import 'environments.libsonnet';
local builders = import 'builders.libsonnet';

local TaskPreparePipeline(name, file, image='prepare-pipeline-container') = builders.Task(name=name, file=file, image=image);
local TaskPerRegion(module, regions, params={}) = concourse.Parallel([builders.Task('%s-%s' % [module, region], module, file=if std.startsWith(module, 'smoketest') then 'test' else 'provision', params=params { REGION: region }) for region in regions]);
local TaskPerArea(module, params={}) = concourse.Parallel([builders.Task('%s-%s' % [module, area], module, params=params { AREA: area }) for area in envs.areas]);
local TaskPerAreaAndRegion(module, regions, params={}) = concourse.Parallel([builders.Task('%s-%s-%s' % [module, area, region], module, params=params { AREA: area, REGION: region }) for area in envs.areas for region in regions]);
local TestPerRegion(module, regions, params={}) = concourse.Parallel([builders.Task('%s-%s' % [module, region], module, file='intranet-conn-test', params=params { REGION: region }) for region in regions]);

local Job(name, dependencies=[], trigger=true, tasks=[], sources=['dev-tools', 'terraform-tools', 'container', 'accounts']) =
  local toA(t) = if std.isArray(t) then t else [t];
  local plan = [concourse.Parallel([builders.Inputs(toA(dependencies), trigger=trigger, sources=sources)]) + toA(tasks);
  concourse.Job(name, plan=plan);
;

local resources = [
  builders.GitResource('source-code', 'products', paths=['modules', product]),
  builders.GitResource('dev-tools', 'dev-tools'),
  builders.GitResource('accounts', 'accounts'),
  builders.GitResource('terraform-tools', 'terraform-tools'),
  builders.DockerResource('container', '$(ACCOUNT).dkr.ecr.eu-central-1.amazonaws.com/$project/public/util-terraform', '0.12'),
  builders.DockerResource('prepare-pipeline-container', '$(ACCOUNT).dkr.ecr.eu-central-1.amazonaws.com/$project/public/util-pipeline', '0.1'),
];

local globalJobs = [
  Job('prepare-pipeline', tasks=[TaskPreparePipeline('validate-product-yaml', 'product-yaml-validation'), TaskPreparePipeline('generate-pipeline', 'pipeline-generation')], sources=['prepare-pipeline-container']),
  Job('functional-area-global', dependencies=['prepare-pipeline'], tasks=[TaskPerArea('functional-area')],
  Job('keycloak', dependencies=['prepare-pipeline'], tasks=[builders.Task('keycloak')])
];

local accounts = [
  { account: 'dev', dependencies: [job.name for job in globalJobs] },
  { account: 'prod', dependencies: ['smoketest-%s' % [env] for env in envs.nonProdEnvironments] },
];

local accountFunctionalArea = [
  Job('functional-area-account-%s' % [account.account], account.dependencies, trigger=envs.continuousDeployment(account.account), tasks=[TaskPerArea('functional-area-access', { ACCOUNT: account.account })]
  for account in accounts
];

local accountNetwork = [
  Job(
    'dispatcher-vpc-%s' % [account.account],
    'functional-area-account-%s' % [account.account],
    tasks=[builders.Task('dispatcher-vpc', 'dispatcher-vpc', params={ ACCOUNT: account.account }),builders.Task('acme-dispatcher-endpoint', 'acme-dispatcher-endpoint', params={ ACCOUNT: account.account })]
  )
  for account in accounts
];

local accountJobs = accountFunctionalArea + accountNetwork;

local EnvironmentJobs(env) =
  local regions = std.objectFields(config.clusters[env]);
  local params = { ENV: env };

  local modules = [
    Job('cluster-base-network-layout-%s' % [env], ['cluster-global-dns', 'functional-area-account-%s' % envs.accountFromEnv(env)] + envs.previousSmokeTest(env), tasks=[TaskPerRegion(module, regions, params) for module in ['domain', 'application-environment-1', 'application-environment-2', 'operations']]),
    Job('customer-resources-vpc-%s' % [env], ['cluster-base-network-layout-%s' % [env]], tasks=[TaskPerRegion(module, regions, params) for module in ['customer-resources-vpc']]),
    Job('cluster-services-%s' % [env], 'cluster-base-network-layout-%s' % [env], tasks=[TaskPerRegion(module, regions, params) for module in ['system-services', 'cluster-backup', 'cluster-management', 'monitoring', 'logging', 'tracing']]),
    Job('functional-area-runtime-%s' % [env], 'cluster-services-%s' % [env], tasks=[TaskPerAreaAndRegion('functional-area-runtime', regions, params)]),
    Job('ingress-%s' % [env], ['cluster-services-%s' % [env], 'dispatcher-vpc-%s' % envs.accountFromEnv(env)], tasks=[TaskPerRegion(module, regions, params) for module in ['private-ingress-app', 'private-ingress-dispatcher', 'access-from-other-accounts']]),
    Job('egress-%s' % [env], ['cluster-base-network-layout-%s' % [env], 'ingress-%s' % [env]], tasks=[
      TaskPerRegion('private-egress', regions, params),
      TaskPerRegion('private-egress-mq-server', regions, params),
      TestPerRegion('intranet-conn-test', regions, (params + {ACCOUNT: envs.accountFromEnv(env)})),
      TaskPerRegion('private-egress-acme', regions, (params + {ACCOUNT: envs.accountFromEnv(env)}))
    ]),
    Job('hello-world-%s' % [env], 'cluster-services-%s' % [env], tasks=[TaskPerRegion('hello-world', regions, params)]),
  ];

  modules + [
    Job('smoketest-%s' % [env], ['%s' % [module] for module in [job.name for job in modules]], tasks=[
      TaskPerRegion('smoketest', regions, params),
    ]),
  ];
;

local nonProdEnvironmentJobs = [{ env: env, jobs: EnvironmentJobs(env) } for env in config.environment_order if std.member(envs.nonProdEnvironments, env)];
local prodEnvironmentJobs = [{ env: env, jobs: EnvironmentJobs(env) } for env in config.environment_order if std.member(envs.prodEnvironments, env)];

local scopedGroups = std.flattenArrays([
  [concourse.Group('global', [job.name for job in globalJobs])],
  [concourse.Group('account-dev', [job.name for job in accountJobs if std.endsWith(job.name, 'dev')])],
  [concourse.Group(jobGroup.env, [job.name for job in jobGroup.jobs] for jobGroup in nonProdEnvironmentJobs),
  [concourse.Group('account-prod', [job.name for job in accountJobs if std.endsWith(job.name, 'prod')])],
  [concourse.Group(jobGroup.env, [job.name for job in jobGroup.jobs] for jobGroup in prodEnvironmentJobs),
]);

local groups = [{
  name: 'all',
  jobs: std.flatMap(function(x) x.jobs, scopedGroups),
}] + scopedGroups;

{
  groups: groups,
  resources: resources,
  jobs: globalJobs + accountJobs + std.flattenArrays([environment.jobs for environment in nonProdEnvironmentJobs + prodEnvironmentJobs]),
}

```

# Scaling up

**Next goal**

**Generate many pipelines**

product\_name: new-product

version: master

clusters:

test:

eu-west-1:

prod:

eu-west-1:

us-east-1:

nodes:

test:

instance\_type: "m5.large"

prod:

instance\_type: "c5.2xlarge"

```
local config = std.extVar('CONFIG')
```

```
# Convert the config to JSON and merge it, so that values are overridden  
CONFIG="$(yq -r '. * .' ../product_defaults.yaml product.yaml \  
    | jq -s 'add')"
```

```
local EnvironmentJobs(env) =  
    local regions = std.objectFields(config.clusters[env]);  
    local params = { ENV: env };  
    local services = ['system-services', 'monitoring', 'logging'];  
  
    local modules = [  
        Job('cluster-services-%s' % [env],  
            'cluster-global-dns',  
            tasks=[TaskPerRegion(module, regions, params)  
                for module in services]),  
    ];  
  
    modules + [  
        Job('smoketest-%s' % [env],  
            ['%s' % [module] for module in [job.name for job in modules]],  
            tasks=[TaskPerRegion('smoketest', regions, params)]),  
    ]
```



**Only works up to a certain point!**

**Some results**

cluster-services-prod #204

started 5h 32m ago

finished 5h 18m ago

duration 14m 18s

+

204203202201200199198197196195194193192191190189188187186185184183182181180179178177176175174173172171170169168167166165164163162161160159158157156155154

↓ source-code

ref 1ac205f3ebc1e06d74f529c952e65442932d784a

↓ dev-tools

ref fce0859c9663b316fe60f5251623988ab4ef3f24

↓ terraform-tools

ref 554b9093eaa01e1a01f5986b558d48479a2d63ee

↓ container

digest sha256:f25a3a7bf3e865ec290b952badebbccddde8ac3475d518f7783835bb969e5d78

↓ accounts

ref 5832e089d5ec31784d8f61039f6abf29603be533

>\_ system-services-ap-northeast-2

>\_ system-services-eu-central-1

>\_ system-services-us-east-1

>\_ cluster-backup-ap-northeast-2

>\_ cluster-backup-eu-central-1

>\_ cluster-backup-us-east-1

>\_ cluster-management-ap-northeast-2

>\_ cluster-management-eu-central-1

>\_ cluster-management-us-east-1

>\_ monitoring-ap-northeast-2

>\_ monitoring-eu-central-1

>\_ monitoring-us-east-1

>\_ logging-ap-northeast-2

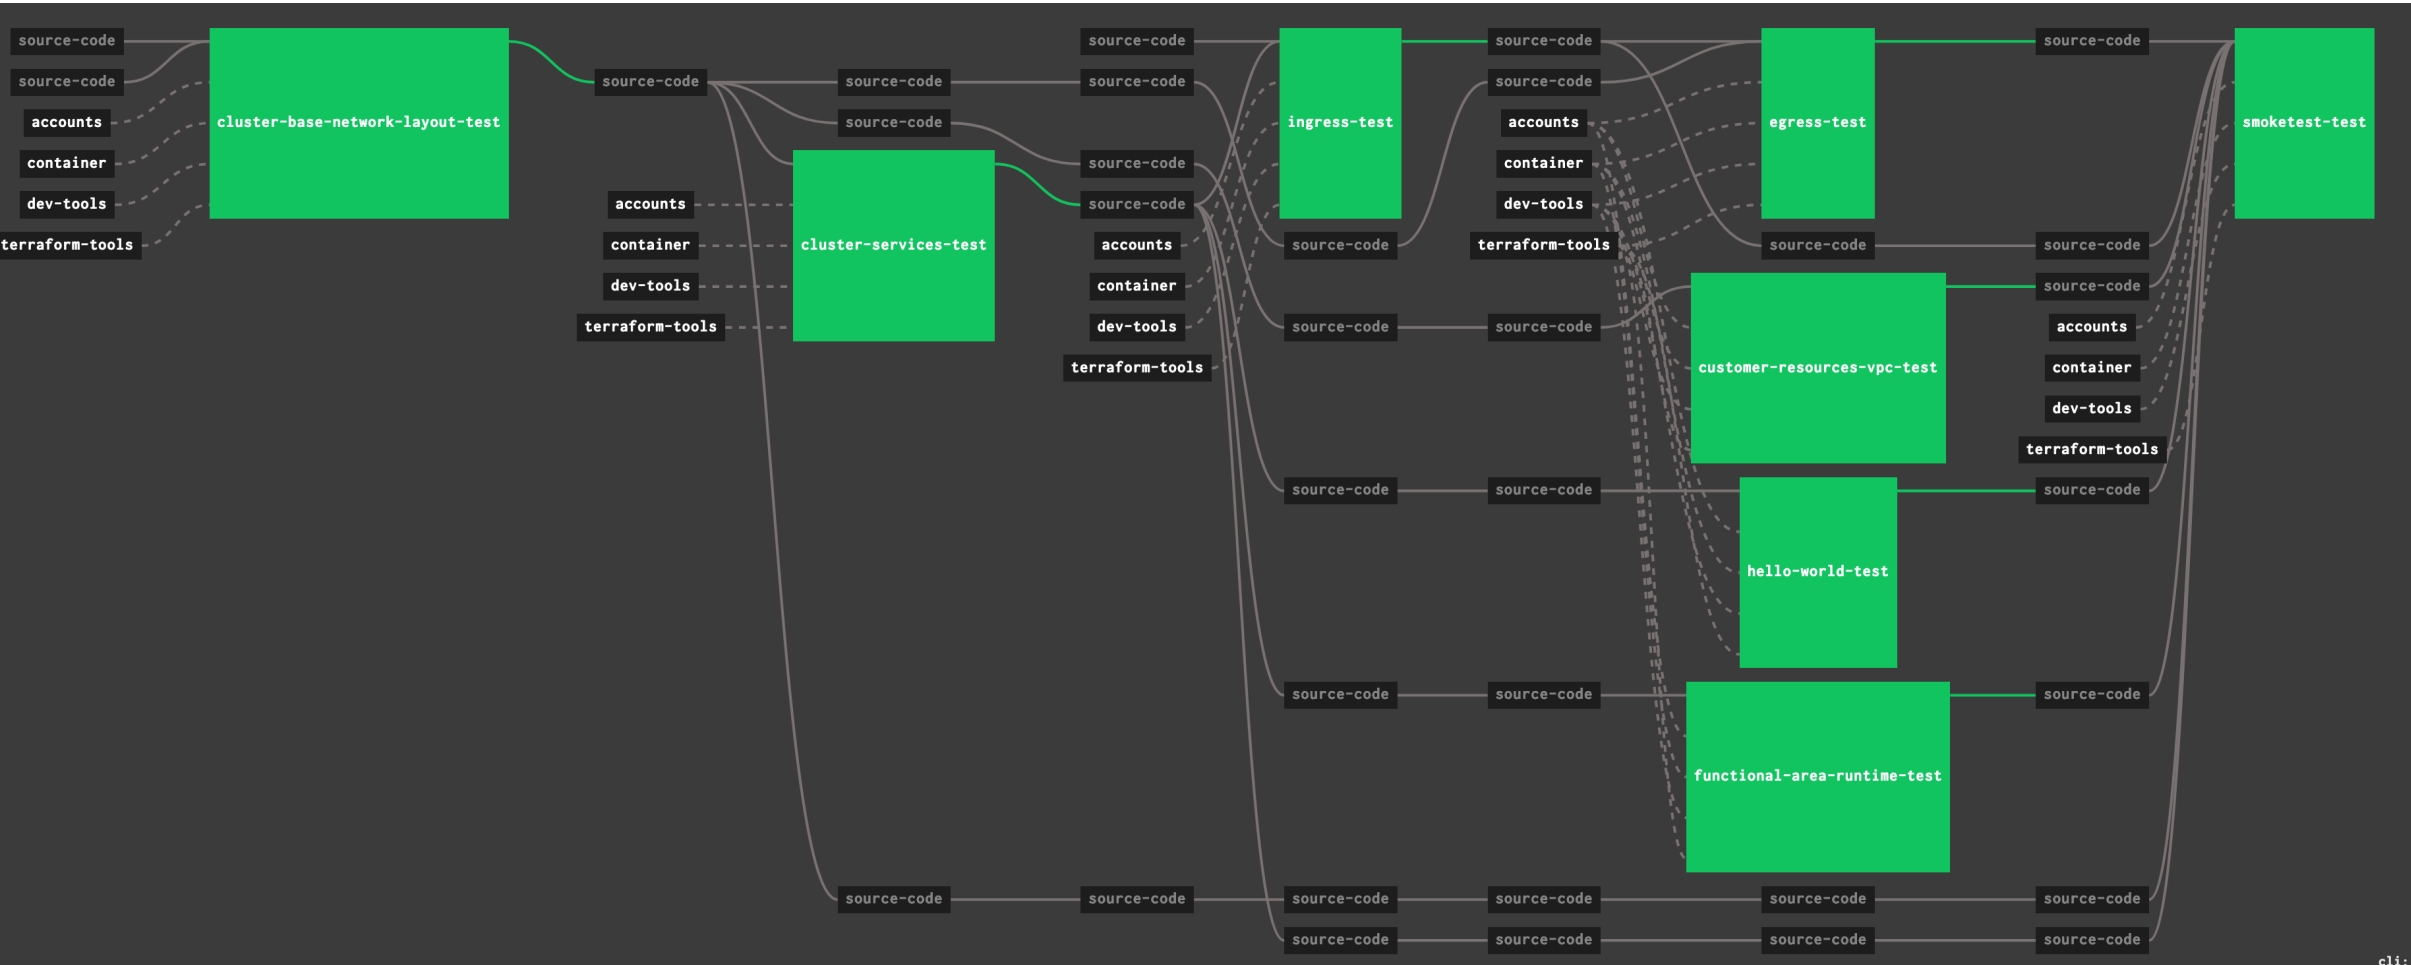
>\_ logging-eu-central-1

>\_ logging-us-east-1


>\_ tracing-ap-northeast-2

>\_ tracing-eu-central-1

>\_ tracing-us-east-1







| Generated     | LOC  |
|---------------|------|
| pipeline.yaml | 7312 |

| Jsonnet                | LOC |
|------------------------|-----|
| pipeline.jsonnet       | 102 |
| concourse.libsonnet    | 54  |
| builders.libsonnet     | 46  |
| environments.libsonnet | 19  |

**[hceris.com/templating-concourse-pipelines-with-jsonnet/](https://hceris.com/templating-concourse-pipelines-with-jsonnet/)**



# Takeaways

**Use Infrastructure as Code**

**Use Infrastructure Pipelines**

**Invest in your tooling**



